

DS 551 Midterm Notes.

Definitions:

- Lec1: Big Data

- 3Vs :
 - Variety : data type variety
 - Volume : magnitude of stored data
 - Velocity : rate of flowing data in and out to users

• Units of measurements:

$$\begin{aligned} \text{- Byte} &= 2^0 \\ \text{- KB} &= 2^{10} \\ \text{- MB} &= 2^10 \cdot 2^{10} = 2^{20} \\ \text{- GB} &= 2^{30} \end{aligned}$$

- Lec2: Firebase, REST and Web API .

- CRUD: - Create - Read - Update - Delete.
 - Create: Post / Put.
 - Read: Get
 - Update: Patch
 - Delete: Delete.

• JSON: JavaScript Object Notation

- value = str | num | obj | array | true | false | null
- object = { string: value }
- array = [values]

- CURL : - PATCH: Requires JSON Obj: ^{inserting} overwrites.
- Put : ~~create~~ / create.

• Filtering Data :

- Order By
 - L : equalTo / startAt / endAt
 - L : limitToFirst / limitToLast
 - L : print

- ordering:
 - b. null → false → true → num → str → obj

Examples:

- Valid Json ?
 - [] ✓
 - {} ✓
 - {[} X
 - [f{}] ✓
 - {"name":john} X
 - {"name": "John"} ✓
 - {"name": 25} X
 - "name" ✓
 - 25.0 ✓
 - {25} X
 - [25] ✓
 - True X
 - true ✓
 - TRUE X
 - Null X
 - "false" ✓

Lec 3 Storage Systems

- Primary : Cache, Memory
- Secondary: Hard disks, SSD
- Tertiary : Tape, Optical Disk

- Access Time: Time taken before drive is ready to transfer data.

- Storage Device Characteristics:

- Capacity
- Cost: Price per byte.
- Bandwidth: how many bytes/sec can be transferred.
- Latency: Time elapsed waiting for response.

- Completion Time = Latency + Size/Bandwidth.

- Workload Completion depends on :

- Workload size (how much work)
- Technology type
- Operation type (read/write)
- Access pattern (random/sequential)

- Access pattern:

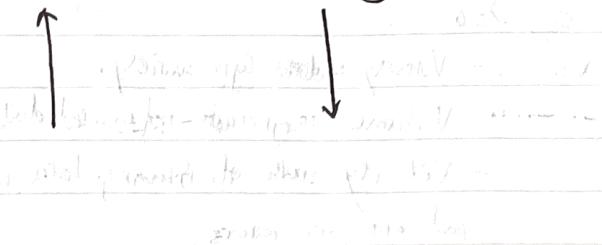
• Sequential: Data to be accessed are located next to each other.

• Random: located randomly on device.

- Magnetic Recording:

- Write head: Applies Electric Current to change direction of mag field.
- Read head: Senses direction

Speed Cost & Energy Consumption



RAM has the highest speed and lowest cost.

SSD has medium speed and medium cost.

Hard disk has medium-low speed and medium-high cost.

Tape has very low speed and very high cost.

SRAM has medium-high speed and high cost.

DRAM has medium speed and medium cost.

Cache has medium-high speed and medium-low cost.

ROM has high speed and very low cost.

RAM has very high speed and low cost.

SSD has medium speed and medium-high cost.

Hard disk has medium-low speed and medium-high cost.

Tape has very low speed and very high cost.

SRAM has medium-high speed and high cost.

DRAM has medium speed and medium cost.

Cache has medium-high speed and medium-low cost.

ROM has high speed and very low cost.

RAM has very high speed and low cost.

SSD has medium speed and medium-high cost.

Hard disk has medium-low speed and medium-high cost.

Tape has very low speed and very high cost.

SRAM has medium-high speed and high cost.

DRAM has medium speed and medium cost.

Cache has medium-high speed and medium-low cost.

ROM has high speed and very low cost.

RAM has very high speed and low cost.

SSD has medium speed and medium-high cost.

Hard disk has medium-low speed and medium-high cost.

Tape has very low speed and very high cost.

SRAM has medium-high speed and high cost.

DRAM has medium speed and medium cost.

Cache has medium-high speed and medium-low cost.

ROM has high speed and very low cost.

RAM has very high speed and low cost.

- Tape.

- Linear Tape: Data recorded on parallel tracks
- Performance Characteristics: high latency, low-cost
- Suitable for low frequency & High access.
- " for high volume data.
- " for sequential access.

HDD size:

- 256 cylinders
- Each cyl has 16 heads
- Each head/platter has 64 tracks
- Each track has 64 sectors
- Sector size: 4 kb

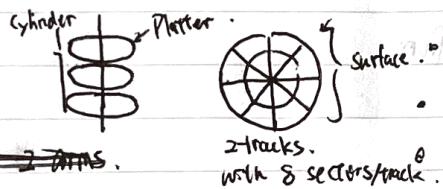
$$256 \times 16 \times 64 \times 4 \text{ kb} = 1 \text{ GB}$$

- File system:

- 2 partitions
 - 1. Metadata, fast read.
 - 2. Blocks of data.
- Directory in XML

- Hard Disk Drive.

- Disk System:



- 1/more magnetic disks
 - each with 2 surfaces.
 - each surface has 2 heads read/write.
 - Each surface has number of tracks.
 - each with number of sectors (blocks)
- ↳ Written all / nothing.

- Completion Time: ~~Seek time~~

- + Seek time: $\frac{1}{2}$ max seek time.
- + Rotation time: $\frac{1}{2}$ full rotation time.

* Transfer time: Data Size / Transmission Bandwidth. Sequential vs. Random:

- Actual Bandwidth = $\frac{|W|}{t}$

- Where $|W|$ size of data.
- t completion time for $|W|$.

$$t = \text{Avg seek time} + T_r + T_t$$

* Random access time of 10 mb

$$t = 2500 \times (T_s + T_r) + T_t$$

- SSD Solid State Drive.

- Performance Characteristics:

- Lower Power Consumption than HDD,
- but more expensive and less capacity
- Limited Lifetime, can only write a limited number of times.
- Better Latency with no T_s and T_r .
- Better on Random access.
- Better on Reading than Writing.

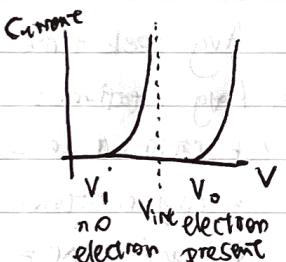


SLC vs. MLC

- ~~File System~~:
- Chip \rightarrow dies \rightarrow planes - blocks - pages - cells.
- cells are made of floating gate transistors.
- Page is smallest to transfer, read, write.
- Block is smallest to be erased.
- SLC:
• Less complex
• Faster - part of SLC is faster than MLC
• More reliable
• Less storage
• More costly

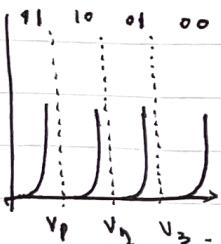
- Hows floating gate works - (SLC)
- Floating gate "floats" inside oxide isolated from other components.
- Electrons are attracted or repelled from FG by applying voltage.
- States of data:
 - 1 : if electron is present.
 - 0 : if not.

- Reading:
• By apply V_{in} , see if current is detected.
1 if detected
0 if not.



- Write & Erase.
 - Write $1 \rightarrow 0$: Applying positive charge will attract electrons.
 - Erase $0 \rightarrow 1$: Applying much higher negative charge will repel electrons. \hookrightarrow which may affect surroundings.

- MLC:



Lec 4 File System.

- Files & Directories
 - Files are stored in blocks
 - Files organized into directories
- CRUD:
 - Create : open(), write()
 - Read : open(), read(), lseek()
 - Update : write(), lseek()
 - Delete : Unlink()

- Inode:

- Created when creating files
- Recording meta-data
- Store locations of blocks & replicas

- Output of stat()

- Access Permissions:

rwx	rwx	r--
$0\ 111\ 111$	100	
$0\ 7\ 7$	4	
- Sector #
- Sector Inode = $(\text{Inode Address} + \text{inode number} \times \text{inode size}) / \text{sector size}$.

Inode: \rightarrow INumber \rightarrow location of inode.

offset INumber = 32.

$$\text{offset Address} = 3 \times 4k + 32 \times 256 \\ = 20k \text{ or } 20480$$

$$\text{Sector \#} = 20k / 512 = 40.$$

Lec 5. HDFS

- Hadoop: Large scale distributed & parallel batch-processing infrastructure.
 - Large scale : Handle large data and computations.
 - Distributed : data distributed among multiple machines.
 - Batch processing : series of jobs without human intervention.
- Distribution.
 - A single namenode storing metadata
 - Directory hierarchy
 - Attributes eg. permission, time, size
 - Mapping to data nodes.

deals with files & metadata. It is local.

- DataNodes: store data after
 - Data are distributed among multiple data nodes.
 - 2 to 3 replicas.
 - Each replica can receive the request.
- Secondary NameNode.
 - Maintain checkpoints of namenode.
 - For recovery

- Block Size: 128 MB.

- Good for fast streaming read of data by sequential pattern.
- Reduce metadata required.

- Client to HDFS.

1. Ask NameNode:
 - Reading: `getBlockLocations()`
 - Writing: `addBlock()`
 2. Ask DataNode:
 - Reading via `readBlock()`
 - Writing via `writeBlock()`
- ~~getBlockLocation()~~:
- Input: File Name,
 - offset (where to start read)
 - Length (how long)
 - Output: Located Blocks.
- (including offset and data node)

Write Block:

- Procedure of writing data to HDFS
 1. To NameNode:
 - `addBlock(ss::src, ss::clientName)`
 - ↳ creates path, permission, client, flag, replicas, size
 - `addBlock(ss::src, ss::clientName)`
 - ↳ Located Block × 1
 2. To DataNode:
 - `block to be written`
 - 1) `writeBlock(block, rest of data nodes, current datanode)`
 - Block break down to packets (4KB)
 - Client send the packets to list of datanodes (Not replicas)
 - Datanodes send packet to Next datanode
 - Acknowledgment → packet removed from ack que (packets)
 - 2) Acknowledge NameNode to update.

lec 6 XML & XPath

Definition: XML = syntax for data

- Benefits:
- Can translate Any Data to XML
 - Over Web
 - Into App

Terminologies:

- Tags: $\langle \text{start} \rangle$ $\langle \text{end} \rangle$
- Elements: $\langle \text{start} \rangle \text{ content} \langle \text{end} \rangle$
- Single Root element
- Attributes: $\langle \text{start} \text{ attr}=\text{"value"} \rangle$
- Must be quoted " or '

DTD: Document Type Definitions

Definition: Markup for describing schema

Well-formed XML: tags are correctly closed

valid: if XML has a DTD and conforms

XPATH:

- //: finding descendants eg. //employee (find all elements with tag)
- [int]: select child by index eg. /root/child[1] (start from 1)
- text(): All text nodes
- *: All element nodes
- nodes(): All nodes
- @: attributes eg. @price
- contains(element, 'text')
- and, or, not()
- |: alternatives

DTD =

```
<!DOCTYPE company [
    <!ELEMENT com element-name (CONTENT) >
    <!-->
]>
```

CONTENT: is character for content

- RE over other elements
- Text only: #PCDATA / #CDATA
- EMPTY
- ANY

RE: Regular Expression

- Optional eg. (firstname?)
- Multiple eg. (phone*)
- Alternation eg. (phone|email)

Lee 7: ER Models

Definitions:

- ER: Entity Relationship
 - Entity: real-world objects described using attributes
 - Attribute: each has an atomic domain
 - Entity set: a collection of similar objects entities
 - Relationship: subset of $A \times B$
- $A = \{a, b, c\}$
 $B = \{1, 2, 3\}$
 $R = \{(a, 1), (b, 2), (c, 3)\}$

What does Arrow means?

- End of the Arrow determines the head of the arrow

Relationships:

- One-to-one: both tables can have only one record on each side of the relationships



- one-to-many/many-to-one:



Only one record in Table A may relate to any number of records in Table B.

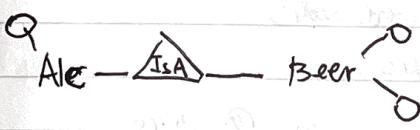


- Many-to-many: multiple records in both tables may be related to any number of records in the other table.



- Binary vs. Multiway: Can convert between.

- Attribute: Attribute of relationship



Subclass:

Definition: Special case, fewer classes, more properties.

ER.	Object-Oriented
- Only one class	- Multiple classes
- All attributes together	- stored separate.

Constraints:

Definitions: Rule that applies to ~~a~~ table all rows of a table or all times

• Keys: one unique attribute ~~as~~ identifier (can consist of multiple attr)

• Single-value constraint: An entity can have at most 1 for a given attr or rel

• Referential Integrity constraints:

referred entity must exist in database exactly once.

• Domain constraints: domain of attr

• General constraints:

(ER) Entity, Entity, Entity

Weak Entity Set:

Definition: the key attr of the set comes from other entity sets.

• Causation:

- Converting Multimany to Binary
- Part-of relationship

Lec 8 Relational Model

Definition: different from ER, everything is represented with tables / relationships

- Domain : Dtype.
- Schema: Relation (att₁, att₂, att_n)
- Instance: Collection of data stored at a moment.

- Tuple
 - element of string × int × string
 - g: t = ("gizmo", 19, "gadgets")
- Relation
 - subset of string × int × string
 - Order is important!

Example Insert t = (123, 'john', 35)

Modification cost	Stability
High	High
Low	Low
Low	Low

function t: A → att domains
 e.g: t(name) = "gizmo"
 A set of tuples / functions
 - Att name is important!
 Insert Employee(123, 'john') = (35)

Transferring ER to Relational :

- Entity Sets: - Set name as relation name.
 - att as att

- Relationships: - Relationship name as relation name.
 - entity att as atts, watch for conflicts

- Many to One: - "Many" Entity as table.
 - "One" as att.

- Many to Many: Not a good idea due to Redundancy

- Weak Entity Sets: No need for additional table for relationship.

- Subclass:

1. OO Approach: Each tree creates a new table. Names in tables are distinct.

2. E/R Approach:
 - A Superclass table
 - Several Subclass Table.
 - Intersects.

3. Null Val Approach: One table

- class as att, classatt as att.
- Null for no class.

eg. Product (name, cat)
 EduProduct (name, cat, ageGroup)

eg. Product (name, cat)
 EduProduct (name, ageGroup)

eg. Product (name, cat, ageGroup).

Lec 9 SQL

DB Modifications

1. Insert:

- Insert into R Values ('a', 'b');
- Insert into R(A) Values ('a');

2. Delete:

- Delete from R where A='a';
- Delete from R where exists (SUBQ);

3. Update:

- Update R set A='a' where B='b';

DDL

1. Create:

- Create Table R (A, B);

2. Drop:

- Drop Table R;

3. Constraints:

- Create Table R (

A CHAR(20),

B VARCHAR(20),

Real / float

INT / TINYINT / SMALLINT ..

4. Key

- Primary

- Unique

5. Default:

- Default 'a'

6. Alter:

- Alter Table R ADD A CHAR(20);

- Alter Table R DROP A;

Lec 10 Constraints

- Keys:
 - Unique and Primary : $\text{create table } R(\text{a int primary key});$
 - Referenced attr must be either - OR $\text{create table } R(\text{a int, primary key(a)});$
 - Null can exist in Unique, not Primary.
- Foreign Key = Referencing relation with Attribute :
 - $\text{create table } R(\text{a int references } B(\text{b})\text{);}$
- Actions taken for Foreign Key Update/Delete:
 - 1. Default : Reject modification.
 - 2. Cascade : Same changes made in R.
 - Delete ~~old~~ R tuple.
 - Update value in R.
 - 3. Set NULL

Policies :

- Cascade ON Delete
- ~~Update~~ Set Null ON Update.

• foreign key (~~a~~) references B(b)
 ON Delete . see null
 ON Update ~~or~~ cascade;

- Value-based constraint : Check $p \leq 50$,
 (at the end) Check 'bar = 'Joe' OR
 $\text{price} \leq 5.00$
- Tuple-based constraint

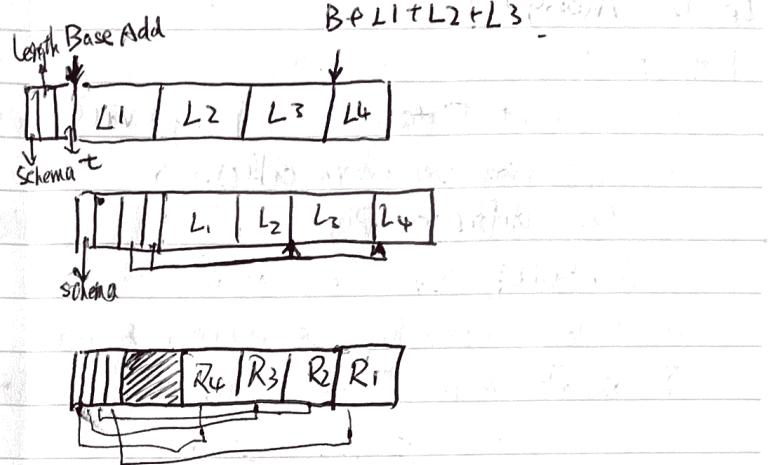
Lec 10 Views

- Virtual Views : Computed on-time
 - Up-to date
 - Slow
- Materialized Views : Precomputed
 - May have stale data
 - Fast

Lec 11

Representation of Data

- Fixed Length
- ~~Variable Length~~ Variable Length
 - Place Fixed Length blocks first records
- Stoaring Records:
 - ~~Records from right~~ from left



External Sorting.

- Buffer Page: Page used to store I/O and intermediate data.

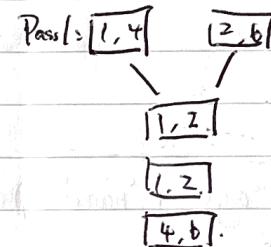
• 2-way Merge Sort

Pass 0: Sort into pairs. (1 buffer)

Pass 1-M: $\frac{[1,4]}{[1,4]} \frac{[2,6]}{[2,6]}$ (3 buffers)

$$\text{Cost} = 2N(\lceil \log_2 N \rceil + 1)$$

$$k = \lceil \log_2 N \rceil + 1$$



~~Cost~~ $\lceil \log_2 N \rceil + 1$

$$\text{eg: } (M=5), B(R)=108.$$

$$P_0: N = \lceil 108/5 \rceil = 22, B = 5.$$

$$P_1: N = \lceil 22/4 \rceil = 6, B = 20.$$

$$P_2: N = \lceil 6/4 \rceil = 2, B = 80.$$

$$P_3: N = \lceil 2/4 \rceil = 1, B = 80.$$

• M-1 way Merge Sort.

$$\text{Cost} = 2B * k$$

$$k = \lceil \log_{M-1}(B/M) \rceil + 1$$

Lec 12. MongoDB

Features:

1. Document Database. $\{ \text{key} : \text{val} \}$ Document \approx JSON Obj.
2. Documents stored in collections
have different types.

3. No SQL

4. Allow hierarchy \approx array & more.

5. Scalable.

$-\text{id}$: unique-key.

Object ID(): 12 byte hex val

- 4-byte: sec since 1970/1/1

- ~~machine~~ 3 bytes = machine.

- 2-byte: process id.

- 3-byte: counter.

- Sorting: `find().sort({ "age": 1, "name": -1 })`

-1 for asc, -1 for desc.

- limit & skip: `find().skip(1).limit(1)`

- Distinct: ~~NO~~ `find().distinct()`

distinct take over find.

- length: `.length`

\Leftarrow

- count: `find().count() / .count()`

- Projections: `find({}).
{ "name": 1, "age": 1 }`

- Renaming: `find({}), { "home": 1, "-id": 0 }
{ address } \Leftarrow`

- Update:

`update({}), { $set: {} }, { multi: true }`

- $\$set$: update instead of overwrite.

- $multi$: multiple,

`updateMany()` \Leftarrow

- Operations:

use `<db>`

show `<db> = db`

`db.dropDataBase()`

`db.createCollection('a')`

`db.a.drop()`

- Adding Documents.

`db.person.insert({ "-id": 1, "name": " " })`

`db.person.insertMany([{}, {}])`

`db.person.find()`

- Operators:

$\cdot \$lt, \$gt, \$lte, \$gte, \$eq, \$ne, \$in, \all

$\cdot \$and, \$or, \$not$.

* $\$and / \or requires [...]

* $\$not$ can not be top level.

eg: ① `find({ $or: [{}, {}] })`

② `find({ "a": { $gt: 25 } })`

- Pattern matching: / Kevin /.

contains \downarrow case insensitive.

`find({ name: { $not: /kevin/ } })`

need: { name: { \$exists: true } } \Downarrow

{ name: { \$not: /kevin/, \$exists: true } }

- On array: $\$elemMatch$.

- Aggregation:

aggregate ({ \$group : { _id: "\$category",
total: { \$sum: "\$colname" } } })

Select category _id, sum(\$total)

From ...

GB. colname

- Having

aggregate (.. , { \$match: { total: { \$gt: 50 } } })

operator: > , <, <=, >=

- Count

aggregate ({ \$group: { ... }, total: { \$sum: 1 } }).

\$sum, \$count, \$size, \$avg

- Multi GB.

aggregate ({ \$group: { _id: { a: "A", b: "B" } } })

\$max, \$min

- Pipeline:

1. \$match → \$group, → \$match → \$sort → \$limit.

- Join

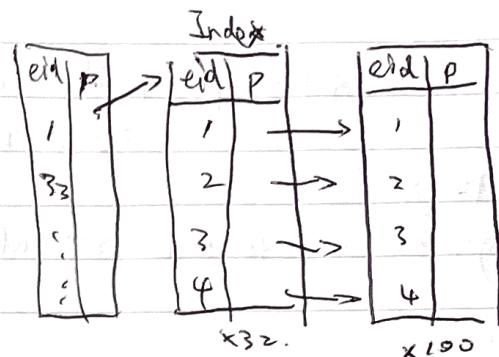
aggregate ({ \$lookup: { from: "A", localField: "a",
foreignField: "b",
as: "c" } })

Lec 13 Indexing

Index is a data structure that speeds up selections on the search key field(s)

- Classification:

- Clustered: records stored in order of search key
- Unclustered: records stored .. not in order
- Search key: subset of fields of a rel.

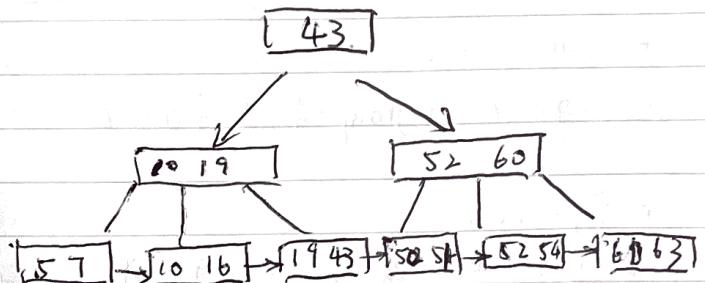


B+ Trees

A search tree that efficiently supports range queries by making leaves into a linked list.

- Rules:

- Order d .
- Except root, every node has $\geq d, \leq 2d$ keys.
- ~~Each leaf has~~



$$\begin{aligned} d &= ? \\ \text{Key} &= 4 \text{ b} \\ \text{Pointer} &= 8 \text{ b} \\ \text{Block} &= 4096 \text{ b.} \end{aligned} \quad \left\{ 2d \times 4 + (2d+1) \times 8 \leq 4096 \text{ b.} \right. \quad \left. d = 170 \right.$$

→ Searching

- 1- Start from root, search left boundary vertically
- 2- locate right boundary while reading all blocks until right boundary (exists or passed).

- Insertion,

- ~~search vertically for the node that needs to be inserted~~
- If overflow ($2d+1$):
 - ↳ split in half, middle to parent.
 - else
 - ↳ stop.

- Deletion

- Search vertically ~
- If $(d-1)$
 - ↳ merge with neighbor.
 - or
 - ↳ borrow from neighbor.

- Splitting

- Parent nodes (internal) don't need to keep node

Lec 14 Query Execution

- Logical Operators: what they do
- Physical Operators: how they do

$\begin{cases} - \text{Scan} \\ - \text{hash} \\ - \text{sort} \\ - \text{index} \end{cases}$

Cost Model

- M : Memory available in main mem.
- $B(R)$: Num of blocks holding R .
- $T(R)$: Num of tuples in R .
- $V(R, a)$: Num of distinct values of att "a" in R .
- Selectivity: diversity of an att in R .
 - clustered: records stored in block R .
 - unclustered: mixed in blocks.

Operator: Scan

cost	clustered: $B(R)$
	unclustered: $T(R)$

One Pass Algorithm

- Selection & Projection
- Cost: $B(R)$
- Duplicate Elim.
- Cost: $B(R)$
- Grouping
- Cost: $B(R)$
- Binary Opr.
- Cost: $B(R) + B(S)$

Nested Loop

- For every $(M-2)$ size R , join with whole table S .
- COST: $B(S) \cdot B(R)(M-2) + B(R)$.
- ~~SKR~~

Two-Pass

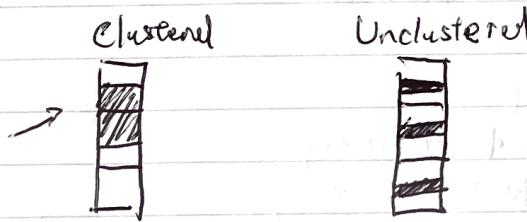
- ~~write~~
- Sorting: DE
- Merge: $B(R)$.
- Grp: same.

Binary Opr.: $3B(R) + 3B(S)$

Join: Sort-Merge Join: $3B(R) + 3B(S)$

* Merge phase cost increase due to extra merge passes.

Simple-Sort Join: $5B(R) + 5B(S)$



Sort-Merge Join

Hashing Based:

- DE: Hash into buckets: $2B(R)$
- Check Duplicate: $B(R)$

Partitioned Hash Join

- COST: $3B(R) + 3B(S)$
- Read S_i , hash it, load one block from R
- Join

Index Based:

- Selection: $T(R) / V(R, a)$

- Join: $B(R) + T(R) B(S) / R$
- $T(R) + T(R) B(T) / V$

- Sort-Merge Join: $B(R) + B(S)$

Lee15. NoSQL

Relational Database

- * Mature and stable
- * Feature-rich versatile: SQL

* ACID:

- Atomicity: All or None opr. executed
- Consistency: DB is consistent with opr.
- Isolation: concurrent opr. don't interfere.
- Durability: can recover from failure.

* Challenges:

- Internet scale Apps.
- Big data, unstructured.
- New workloads

Consistency:

- Eventual Consistency:

- Weak Consistency
- Update will eventually show up.

- Inconsistency Window:

Time between update acknowledged to

User / all replicas

Determined by Communication Delay

Lead on sys.

Number of replicas

CAP: Can have at most 2.

- Consistency: Every read receives most recent write/error.
- Availability: Every request receives a response
- Tolerance to Net Partitions: Sys continue to operate when messages are dropped.

NoSQL

NoSQL Database = Not only SQL

- * Flexible (non-relational)

- * Scalable horizontally

- * Data replicated

- * Weaker consistency

- * High availability

eg:

- Firebase

- MongoDB

- DynamoDB

DynamoDB:

- Schema-less (other than primary key)

- DB structure.

DB

↳ Table

↳ Items / rows

↳ Attributes

- Partition key: Wide range of value

- Sort key: searching within a partition

- Datatypes:

- String, Binary, Number, Stringset

- Numberset, Binaryset, Map,

- List, Boolean, Null

Lec 1b. MapReduce

MapReduce : A computational paradigm

- Parallel

- Distributed

A MapReduce Job consists of a number:

- Map tasks = Data Transformation
- Reduce tasks = Combines results
- Internal Shuffle-tasks = distribute tasks.

Mapper:

- ▷ Invokes Map func for each input key-value pair.

Reducer:

- ▷ Invokes Reduce func for each different intermediate key.

Shuffler:

- ▷ All intermediate key-value pairs are sent to the same Reduce Task.
- Partitioner = Hashing / ..

FileInputFormat:

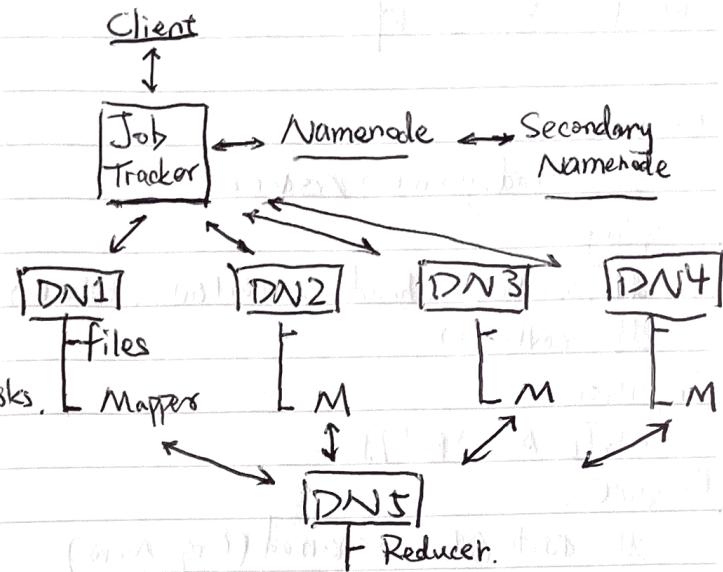
- Take paths to files
- Read all files in paths
- Divide files into Input Splits.
 - ↳ One Map Task for each Split.
- Subclass: InputFormat
 - ↳ Subclass: RecordReader.

FileOutputFormat:

- Output from Reducer.
- Each write in separate part-nnnnn.
 - ↳ n = Reducer out
 - ↳ nnnnn = partition #

Combiner:

- Require = Operation be Associative & Commutative.
- Input & Output Format = Map Out & Reduce In.



1. Client request to Job Tracker.
2. Job tracker:
 - ▷ Ask NameNode for loc.
 - ▷ Assign task to Task Tracker at Data Node.
 - ▷ Reassign if fails.
3. Task tracker:
 - ▷ Signal back to Job Tracker.
 - ▷ Monitor status of task.
4. Mapper: $\langle k, v \rangle \rightarrow \langle k', v' \rangle$
5. Shuffler:
 - ▷ Partition Map output in buffer into R (Reducers Tasks).
 - ▷ Sort Partitions and write disk.
 - ▷ Merge on disk.
 - ▷ Notify Job Tracker
 - ▷ Copy data from all Mappers
 - ▷ Merge data
 - ▷ ...

Lec 17 Spark DF

Creation:

`spark.read.json() / csv()`

`spark.createDataFrame([(1, 1), (2, 2)], [colA, colB])`

Display:

`df.show(n) / .head(n) / .tail(n) / .take(n)`

`df.collect()`

Projection:

`df[['A', 'B']]`

Rename:

`df.withColumnRenamed('Org', 'New')`

Add Column:

`df.withColumn('colName', val)`

Filter:

`df[(A > a) & (B > b)]`

Distinct:

`df.dropDuplicates() / df.distinct()`

Aggregate ~~with~~ ~~for~~:

`df.agg({'ColA': 'max'})`

`df.agg(F.max('ColA').alias('maxColA'))`

GroupBy:

`df.groupBy(['colA', 'colB']).max()`

OrderBy:

`df.orderBy(['colA', 'colB'], ascending=[T, F])`

Agg Func:

`count, max, min, avg/mean, sum`

Limit: `limit(n)`

Join:

`df1.join(db2, (db1.a == db2.a) & (db1.b == db2.b))`

`df1.join(db2, 'colA')`

Intersect: `df.intersectAll() / df.intersection()` # removes duplicates

Subtract: `df.subtract(db2)`

Lec 18 Spark RDD

Motivation:

- Hadoop ill-suited for
 - ↳ Reuse looping data
 - ↳ Interactive exploration.

RDD: Resilient Distributed Dataset

- Read only, partitioned collection
- Parallel Operate on partitions
- Fault Tolerant

Create Spark context

`sc = SparkContext()`

Create:

`sc.textFile()`

`sc.parallelize(list, num)`

Reduce:

`data.reduce(func) \Rightarrow val` Reduce By key
 $\#$ func commutative & associative. `data.reduceByKey(func) \Rightarrow RDD`.
 $\#$ data = $[(k, v)]$

Map:

`data.map(func)`

Filter:

`data.filter(func) # boolean func`

Groupby Key

`data.groupByKey() \Rightarrow [(k, [v, ...]), ...]`

map Value:

`data.mapValue(func) \Rightarrow [(k, f(v)), ...]`